

To ask questions on Meetup topic, join at:

[#EmbeddMeet](https://slido.com)

Za spoluprácu na dnešnom
Meetupe ďakujeme:

Matúš Chochlík (Speaker)

Za organizáciu a propagáciu
eventu ďakujeme:



FAKULTA RIADENIA A INFORMATIKY
ŽILINSKÁ UNIVERZITA V ŽILINE



ŽILINSKÁ UNIVERZITA V ŽILINE
Elektrotechnická fakulta

GlobalLogic

robíme **it**

The poster features a teal background with a faint circuit pattern. At the top, there is a stylized logo consisting of three horizontal black bars on the left and two vertical yellow bars on the right, with a yellow circle in the center of the right vertical bar. Below this, the text "Embedd" is in large black font, and "Meet" is in large yellow font. Underneath, "Zilina" is written in black font inside a yellow rectangular box. To the right of the box are three horizontal black bars. At the bottom, the text "FUNPAGE" is in yellow, and the Facebook URL "www.facebook.com/EmbeddMeetZilina" is in white.

**Embedd
+++ Meet**

Zilina ---

FUNPAGE
www.facebook.com/EmbeddMeetZilina

External process synchronization for fun and profit

November 19, 2019

The LLVM project

- “a collection of modular and reusable compiler and toolchain technologies”
 - `llvm` – code generator and optimizer for many CPUs
 - `lldb` – debugger
 - `lld` – linker
 - `clang` – C/C++/Objective-C compiler
 - `clang-tidy` – C/C++/Objective-C static analysis tool
 - `clang-format` – C/C++/Objective-C code format tool
 - ...
- <https://llvm.org/>
- <https://github.com/llvm/llvm-project>

Let's build `llvm` and `clang`

```
$ git clone https://github.com/llvm/llvm-project.git
```

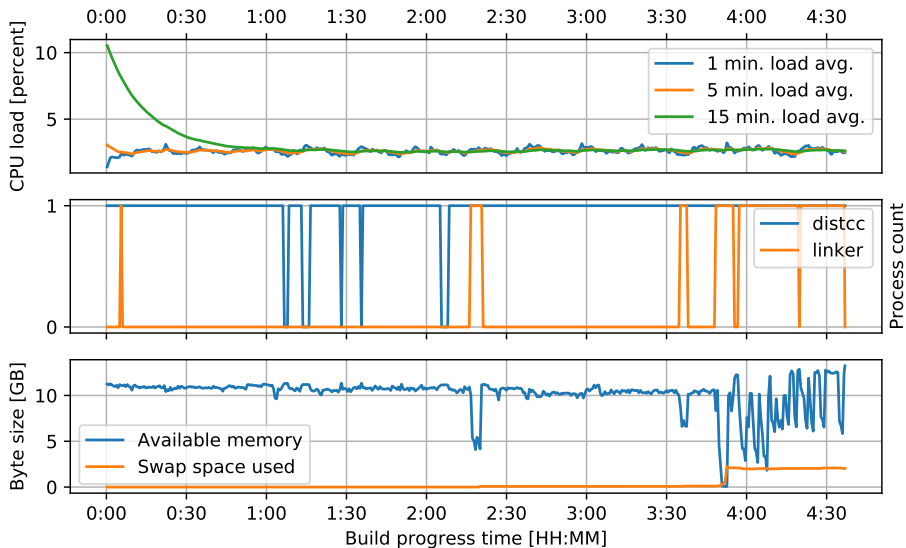
```
$ mkdir _build
```

```
$ cd _build
```

```
$ cmake -DCMAKE_BUILD_TYPE=Debug ..
```

```
$ make
```

make – system resource usage



That didn't go so well...

- System resource usage is low
- Build takes cca. 5 hours – unacceptably long
- Let's try some tricks:
 - use `ccache`
 - use more `make` jobs
 - use `distcc` for distributed compilation

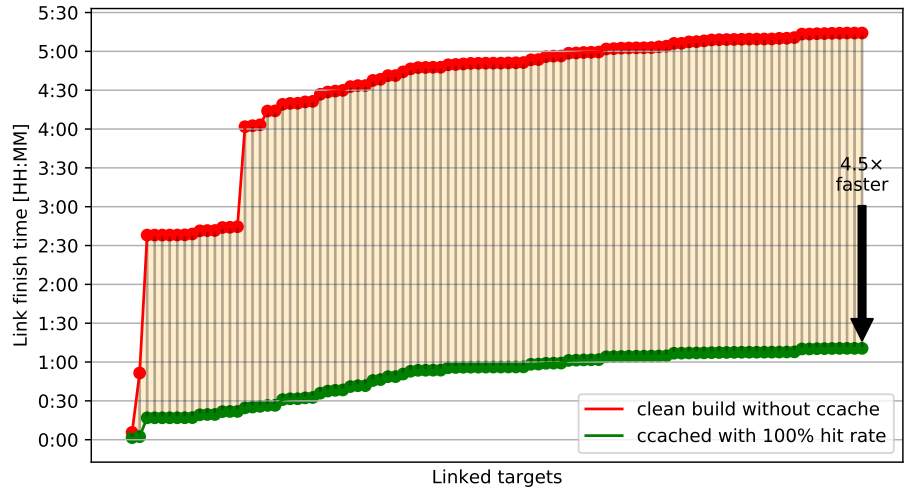
ccache – compiler cache

- “speeds up recompilation by caching previous compilations and detecting when the same compilation is being done again”
- <https://ccache.dev/>


distcc – distributed compilation for C and C++

- “distributes compilation of C or C++ code across several machines on a network”
- <https://github.com/distcc/distcc>

make – uncached clean build vs. 100% cached



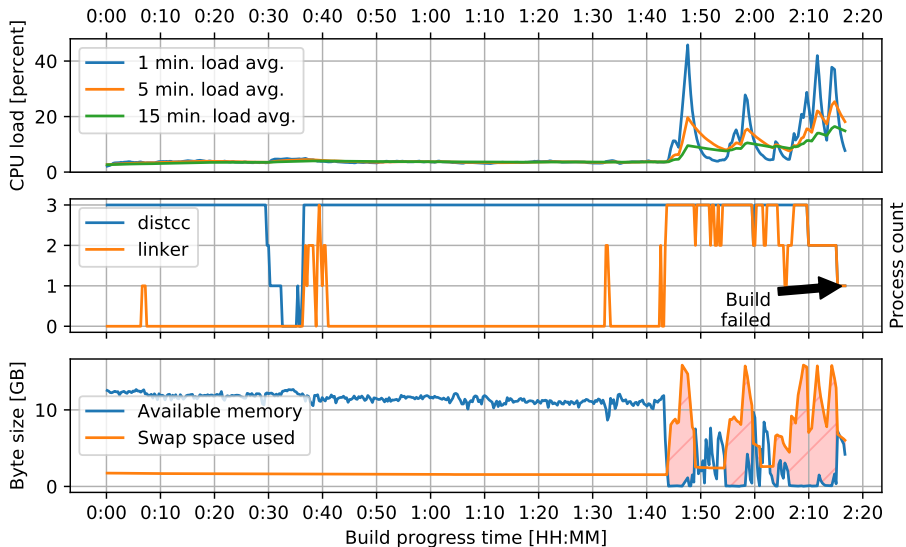
More jobs – make clean && make -j 2



That's somewhat better...

- System resource usage is still low
- Build takes more than 2 and half hours – still too long
- Parallelization shows potential

Keep going – `make clean && make -j 3`




What happened?

- System usage starts to climb
 - especially memory usage
 - signs of correlation with linker execution
- Build *failed* after 2 hours 20 minutes – looooong
 - Linux OOM¹ process killer, kills some linker jobs
- Parallelization still shows potential

¹Out Of Memory

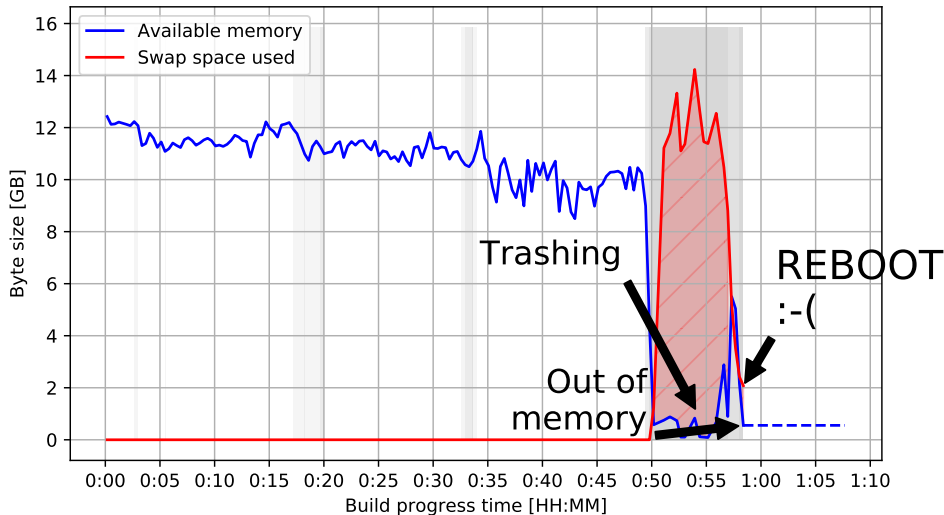
Meh. MOAR jobs! `make clean && make -j 7`



Not great, not terrible

- System usage very high toward the end of the build
 - especially memory usage
 - still correlated with linker execution
- Build *failed* after 1 hours 5 minutes – getting better
 - OOM process killer, kills some linker jobs
- Parallelization rules!

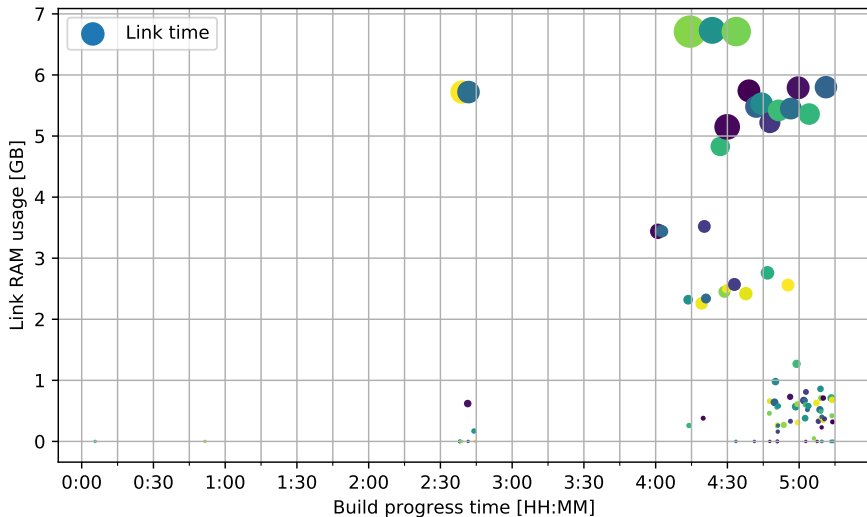
Still come cores left... `make clean && make -j 9`



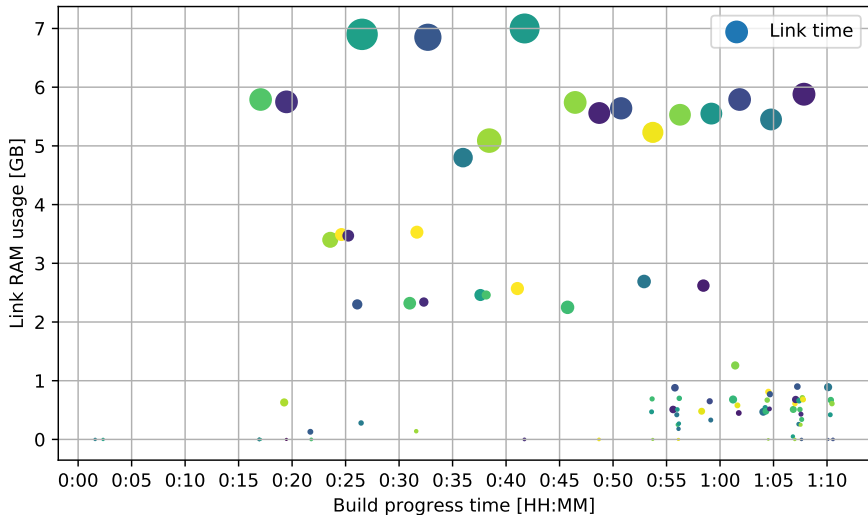
What have we learned so far

- Most of the build time is spent by compilation
- Compilation caching – good.
- Compilation paralelization – good.
 - We have not even really used `distcc` yet.
- Linking paralelization – bad! Why?
 - Let's have a closer look!


Linker memory usage vs link time – no cache



Linker memory usage vs link time – cached



Link target memory requirements – classification




Linking `llvm` and `clang` in a nutshell

- Around 100 different link targets
 - Executables, Shared libraries
 - Many of them are small.
 - Few of the **huge**.
- Most of the linking is done towards the end of the build process.
 - ⇒ Using many parallel jobs, multiple instances of linker run concurrently.
 - ⇒ Many big targets are linked at the same time.
 - ⇒ **OOM!**
- BTW: so much swapping is bad for the SSDs.

Making parallel linking work

- Let's try some additional tricks:
 - Use GNU `gold` instead of GNU `ld`.
 - Can link only in ELF format.
 - + Faster than `ld`.
 - + Uses less memory than `ld`.
 - Use `zram`.
 - a Linux kernel module for creating a compressed block device in RAM
 - can be used as swap device – part of swap is in compressed RAM
 - Make only some targets
 - Practical when developing and debugging.
 - Sometimes you just need to do `make all`.

```
zram + distcc + ccache + gold + make -j 12
```



If only...

- ... we could
 - prevent so many big targets from being linked at once,
 - prevent excessive swapping to disk,
- but still
 - use the hardware resources efficiently,
 - use available remote cores via `distcc` to full extent.

Let's synchronize the execution of `ld` / `gold`

- `atmost` to the rescue!
 - Simple utility written in C.
 - Takes a command-line (executable + arguments).
 - Waits until specified conditions are met.
 - Executes the command-line.
 - <https://github.com/matus-chochlik/various/atmost>

atmost – description

“The atmost utility allows to limit concurrent execution of a single executable or a set of executables, depending on various criteria like CPU load, memory and swap usage, thermal zone temperatures, current number of I/O operations, etc., or just by a maximum number.”

RTFM – `man atmost`

ATMOST(1)

General Commands Manual

ATMOST(1)

NAME

`atmost` - tool for limiting the concurrent execution of a specified executable.

SYNOPSIS

```
atmost [-v|--verbose] [-f|--file file-path]
        [-s|--socket socket-path] [-r|--reset]
        [-i|--sleep-interval seconds] [-c|--print-current]
        [-C|--print-all-current] [-n|--max-instances count]
        [-l|--max-cpu-load-1m percent]
        [-L|--max-cpu-load-5m percent] [-m|--min-avail-ram percent]
        [-M|--min-free-ram percent] [-S|--min-free-swap percent]
        [-p|--max-total-procs count] [-tc|--max-cpu-temp temp]
        [-tg|--max-gpu-temp temp] [-tb|--max-bat-temp temp]
        [-io|--max-io-ops count] [-nw|--max-nw-speed speed]
        [-snw|--slow-nw-speed speed] [-- executable [args...]]
```

```
atmost --help
```

`atmost -n NUMBER -- EXECUTABLE` – basic usage

- Limits the concurrent execution of `EXECUTABLE` to the specified `NUMBER` of instances.
- Uses an IPC semaphore set².
 - *Acquires*³ the semaphore before executing `EXECUTABLE`.
 - *Releases* the semaphore after executing `EXECUTABLE`.
 - `realpath`⁴ of `EXECUTABLE` is used as the semaphore set token⁵.

²see `man 2 semget`

³see `man 2 semop`

⁴see `man 3 realpath`

⁵see `man 3 ftok`

Using `atmost` – wrapper scripts

- In order to synchronize an `EXECUTABLE` with `atmost` :
 - Create a shell script with the name of `EXECUTABLE` .
 - From the script call `atmost` with appropriate arguments.
 - Put the wrapper script to `${PATH}`

```
$ vim /opt/bin/ld
```

```
#!/bin/bash  
atmost -n 2 -- /usr/bin/gold "${@}"
```

```
$ chmod u+x /opt/bin/ld
```

```
$ export PATH="/opt/bin:${PATH}"
```

```
atmost -n 2 -- ld
```

- + Empirically we have found that 2 instances of `ld` can run safely with 16GB RAM.
- + Simple
- + Low overhead
- Too coarse and restrictive
 - The majority of the `llvm` targets has low RAM requirements for linking.
 - Targets from other projects also have low linker RAM requirements.
 - We could safely use more parallel link jobs.

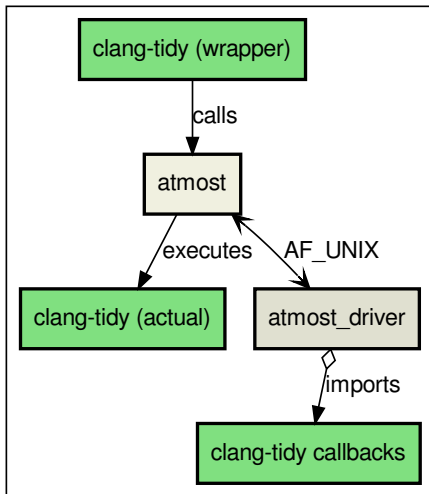
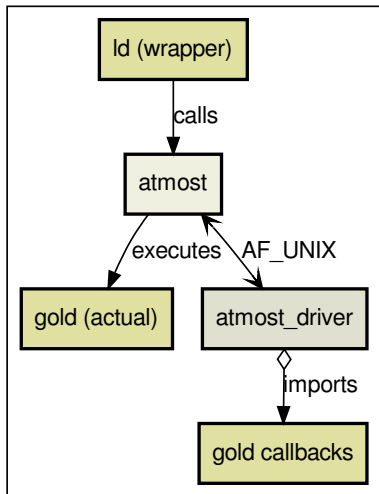
What if...

- ...we could **determine** if it is safe to start the linker per each invocation?
- Enter **atmost_driver**!
 - Running as a local **AF_UNIX** socket server.
 - **atmost** is the client (with the **--socket PATH** option).
 - When **atmost** is started, it connects to the *driver* and sends information about the wrapped **EXECUTABLE**.
 - When the *driver* determines that it's safe to run **EXECUTABLE**, it responds to **atmost**.
 - Then **atmost** executes **EXECUTABLE**.

Let's make it even more flexible

- `atmost_driver` callbacks:
 - The `atmost_driver` implements the reusable common server functionality.
 - The specific logic determining when it is safe to start executing, is implemented in a set of callbacks.
 - `load_callback_data`
 - `save_callback_data`
 - `process_initialized`
 - `let_process_go`
 - `process_finished`

atmost + atmost_driver + callbacks



Callback – `load_callback_data`

```
def load_callback_data():  
    resources = do_load_resources(...)  
    return resources
```

- Called once, when the `atmost_driver` is being started.
- Typically does startup initialization.
- Can load resources used by the other callbacks.

Callback – process_initialized

```
def process_initialized(resources, proc):
    # handle new process
    proc.set_callback_data(callback_data)
```

- Called when `atmost` sends the driver information about a new process.
- The `proc` parameter provides information about the new process.
 - PID, command-line, environment, working-directory, etc.
 - The callback can store its bookkeeping data in `proc`.

Callback – `let_process_go`

```
def let_process_go(resources, procs):  
    if can_process_start(resource, procs):  
        return True  
    return False
```

- Called repeatedly.
- Determines if a process can start executing the synchronized executable.
- The `procs` argument contains info about *all* currently managed processes, split into 3 groups.
 - `active` – processes that are already executing.
 - `waiting` – processes that are waiting for execution.
 - `current` – the process to be let go.

Callback – `process_finished`

```
def process_finished(resources, proc):  
    callback_data = proc.callback_data()  
    # cleanup info about the process
```

- Called when a process handled by the driver has finished.
- The callback can retrieve its bookkeeping data from `proc`.

Callback – `save_callback_data`

```
def save_callback_data(resources):  
    do_cleanup_resources(resources)
```

- Called once, when the `atmost_driver` is being shutdown.
- The `resource` parameter is the return value from `load_callback_data`.
- Can cleanup and/or save resources used and data generated during the run.

RTFM – `man atmost_driver`

ATMOST(1)

General Commands Manual

ATMOST(1)

NAME

`atmost_driver` - driver server for the `atmost` command.

SYNOPSIS

`atmost_driver [-s|--socket socket-path] [-c|--callbacks file-path]
[-u|--update interval]`

`atmost_driver -h`

`atmost_driver --help`

Back to linking...

- What if the memory usage can be determined from the linker command-line arguments?
 - 1) Run a lot of builds of various projects.
 - 2) Gather a lot of data
 - command line arguments,
 - input file sizes,
 - actual linker memory usage,
 - etc.
 - 3) ???
 - 4) Profit!

What exactly is step 3) ???

- Feed the gathered data into a machine learning model and train it.
- Use the trained ML model to **predict linker memory usage** from command-line arguments.
- Integrate the ML classifier into a `atmost_driver` callback script.
 - Keep track of available memory.
 - Predict the memory usage of new instances of linker.
 - Let the new linkers go only if there is enough available memory.


The ML part

- Uses the `neural_network.MLPClassifier` from the `sklearn`⁶ Python package.
- Inputs:
 - optimization level, the PIE⁷ flag,
 - count and combined size of input shared and static object files,
 - etc.
- Output:
 - the predicted memory requirements, as a multiple of 256MB chunks,
 - i.e. if the output is 5 the predicted size is $5 * 256[MB] = 1.25[GB]$.

⁶<https://scikit-learn.org/stable/documentation.html>

⁷position-independent executable


Classifier prediction accuracy



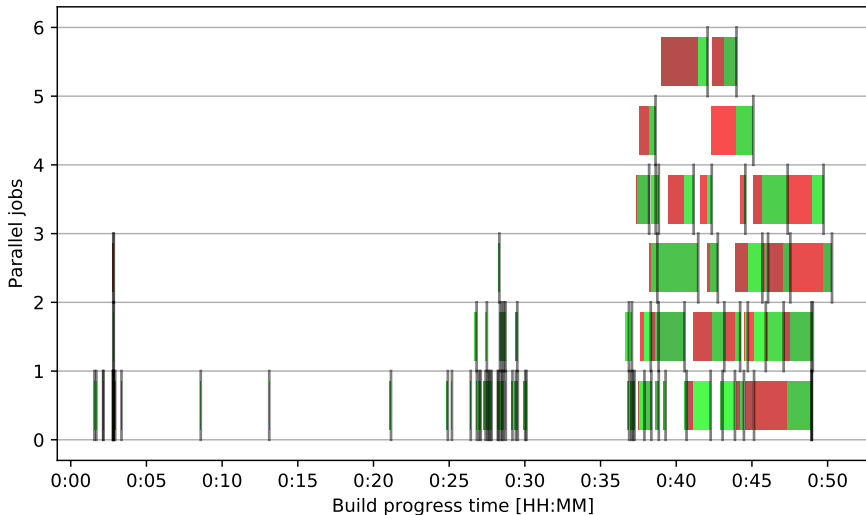
Integrating classifier with `atmost_driver`

- `load_callback_data` – load the trained ML model.
- `process_initialized` – use the model to predict memory usage.
- `let_process_go` – only if predicted linker memory usage is less than current available memory.
- `process_finished` – output actual memory usage that can be stored and used for additional model training.


Linker schedule (no ccache, 16GB RAM) – `make -j 10`




Linker schedule (no ccache, 16GB RAM) – `make -j 20`




Linker schedule (no ccache, 8GB RAM) – `make -j 10`



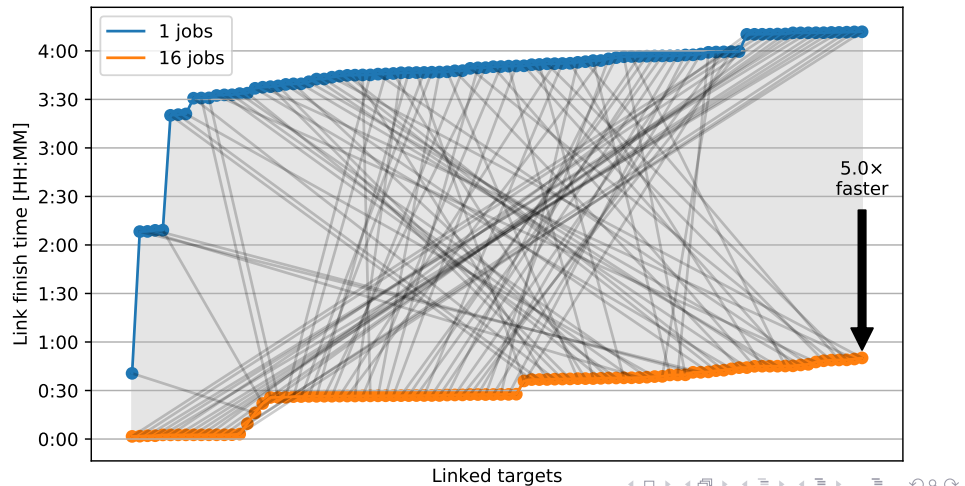
Linker schedule (no ccache, 8GB RAM) – `make -j 20`



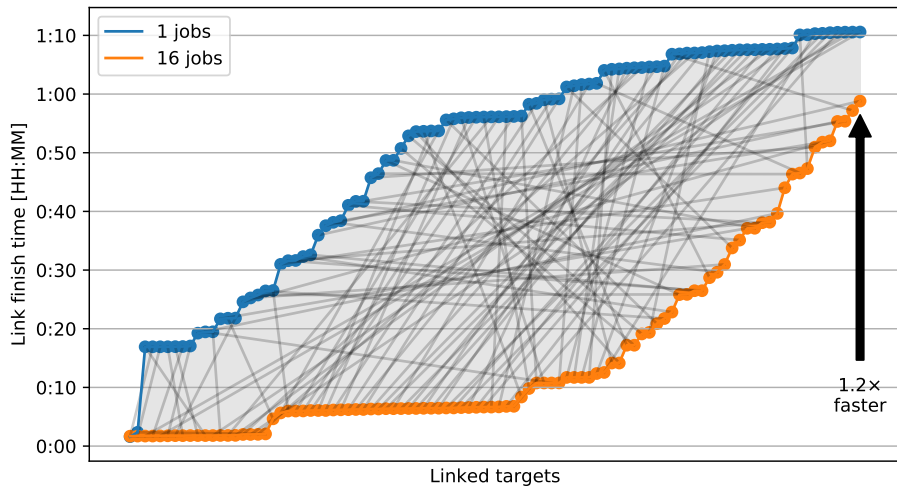
Linker schedule (ccached, 8GB RAM) – `make -j 20`



Link order vs. time (no ccache, 16GB, 4+12 cores, 16 jobs)




Link order vs. time (ccached, 8GB, 8+8 cores, 16 jobs)




Parallelization statistic indicators

- *Run time with j jobs*
 - T_j
- *Speedup with j jobs*
 - $S_j = T_1 / T_j$
- *Marginal speedup with j jobs*
 - $M_j = T_{j-1} / T_j$
- *Efficiency with j jobs*
 - $E_j = S_j / j$

Parallelization statistics (no ccache, 8GB, 8+8 cores)



Parallelization statistics (ccached, 8GB, 8+8 cores)



Conclusions

- `atmost` is a highly flexible tool for synchronizing execution of parallel processes,
 - improves build times by allowing many parallel `make` jobs while serializing huge link jobs,
 - prevents system stalls, crashing and reboots due to low memory conditions,
 - makes building big projects safer,
 - allows to build `llvm` on systems with low RAM capacity,
 - reduces the amount of writes to swap disk partitions.

the *fun* part

- Learned new things:
 - new Python 3 features
 - Python *SciKit-Learn* package
 - Python *MatPlotLib* package
 - Secure `distcc` setup over `ssh`.
 - etc.

the *profit* part

- Stable and relatively quick `llvm` / `clang` builds.
- Reasonable HW usage.
- Saving the SSDs from excessive wear.
- Not getting random reboots during development and building.

the *moral* of the story

Go get more RAM!⁸

⁸if you can

Thank you!

Questions?

`https://github.com/matus-chochlik/
various/atmost/presentation/embeddmeet.pdf`

Extras

atmost – additional options

- Besides the basic semaphore-based synchronization and the customizable driver-based synchronization, **atmost** supports synchronization depending on:
 - 1 and 5 minutes average load,
 - amount of available RAM,
 - amount of free swap,
 - CPU temperature,
 - network speed,
 - battery status,
 - etc.

atmost – average load

- -l or --max-cpu-load-1m
- -L or --max-cpu-load-5m

Start executing only if 1 minute average load is less than or equal 15%:

```
atmost -l 15 -- executable arguments...
```

Start executing only if 5 minute average load is less than or equal 8.5%:

```
atmost -L 8.5 -- executable arguments...
```

atmost – memory / swap usage

- `-m` or `--min-avail-ram`
- `-M` or `--min-free-ram`
- `-S` or `--min-free-swap`

Start executing only if at least 20% of RAM is available:

```
atmost -m 20 -- executable arguments...
```

Start executing only if at least 70% of swap is free:

```
atmost -S 70 -- executable arguments...
```

atmost – total process count

- `-p` or `--max-total-procs`

Start executing only if the number of currently running processes is ≤ 1000 :

```
atmost -p 1000 -- executable arguments...
```

atmost – thermal zone temperatures

- `-tc` or `--max-cpu-temp`
- `-tg` or `--max-gpu-temp`
- `-tb` or `--max-bat-temp`

Start executing only if the CPU temperature is less than or equal to 70°C:

```
atmost -tc 70 -- executable arguments...
```

Start executing only if the battery temperature is less than or equal to 55.5°C:

```
atmost -tb 55.5 -- executable arguments\  
...
```

atmost – modifiers

- Modifiers allow to change the limits by a specified amount under special conditions.
 - when running on battery,
 - when connected to “slow” network,
 - when disconnected from network.

atmost – on A/C vs. on battery

- `-batt` or `--on-battery`

When on A/C, start only if 1 minute average load is less than or equal to 15%,
when on battery, start only if average load is less than or equal to 5% = (15 - 10):

```
atmost -l 15 -batt 10 -- executable \
arguments...
```

atmost – combining limits and modifiers

When on A/C, start only if 1 minute load is $\leq 20\%$ and 5 minutes load is $\leq 10\%$,
when on battery, start only if 1 minute load is $\leq 8\%$ and 5 minutes load is $\leq 4\%$:

```
atmost -l 20 -batt 12 -L 10 -batt 6 -- \
executable ...
```

Limiting concurrent instances of clang-tidy

- Is it useful to limit concurrent execution of clang-tidy ?
 - 8 core CPU,
 - 2 local cores for compiling,
 - 4 local cores for compiling,
 - 6 local cores for compiling,
 - the rest used by clang-tidy .
- Measured on the test-suite of the OGLplus⁹ project.

⁹<http://oglplus.org/>


```
atmost -n  $\{N\}$  clang-tidy " $\{@\}$ "
```

